

3rd Assignment

January 26, 2017

1 Group Members

Pierre Mercatoris – Pablo Bordons Estrada - Sergio Gámez Ruiz de Olano – Mohammadmehdi Fayazbakhsh

```
In [72]: %matplotlib notebook
         from __future__ import division
         import numpy as np
         import time

         nsample = 1000
         nvariables = 100
         X0=np.ones([nsample, 1])
         X1=np.random.uniform(0, 10, ([nsample,nvariables]))
         X=np.concatenate([X0, X1], axis=1)
         error=np.random.normal(0, 1, (nsample,1)) # Normal random error
         beta=np.random.uniform(-5, 5, size=([nvariables+1, 1]))
         Y=np.dot(X, beta)+error
         print(beta)

[[ 1.28069035]
 [ 2.12055007]
 [-3.35364129]
 [ 2.49180722]
 [-2.8342821 ]
 [ 2.01194948]
 [-2.85385757]
 [ 0.3348794 ]
 [ 1.54505771]
 [ 3.76039123]
 [-3.08866171]
 [-0.51838654]
 [ 4.079594  ]
 [ 3.20824974]
 [ 4.02375764]
 [ 1.82350012]
 [ 2.73740613]
```

[4.83990623]
[4.78336779]
[-2.28676607]
[-2.59697585]
[-4.13702962]
[-2.42724456]
[4.51397344]
[-4.76847074]
[4.55519664]
[4.76617981]
[-0.63032856]
[-4.11676195]
[1.24741552]
[-3.9683488]
[-1.08071633]
[-0.25369959]
[-2.47735284]
[2.05335703]
[4.89243931]
[2.84978551]
[0.31243966]
[4.83657486]
[0.43747194]
[-0.77572563]
[3.27644737]
[-1.55841983]
[0.09359911]
[2.1735441]
[1.83173806]
[4.21374777]
[2.44924523]
[-0.69330504]
[-1.93368367]
[4.97225515]
[2.40465098]
[-2.23907213]
[4.66306122]
[-2.9983419]
[-2.31792161]
[-4.7092036]
[2.54629324]
[3.39634]
[4.70833798]
[4.21943738]
[-3.46481252]
[-4.70657762]
[-0.37296535]
[0.14072207]

```

[ 4.26372857]
[ 4.68691278]
[ 1.20068953]
[-4.92552407]
[ 0.61007361]
[ 3.98499229]
[ 2.98152954]
[ 4.09879831]
[ 4.92827865]
[-2.10621153]
[ 3.96988303]
[ 1.05428171]
[-1.58649877]
[ 2.41262391]
[-4.68382943]
[-3.00209887]
[ 2.17824054]
[ 4.62864372]
[-3.42174215]
[ 3.24110498]
[-4.59041032]
[-0.10349687]
[ 4.74605444]
[-4.81932749]
[ 2.84527302]
[-2.46052222]
[ 0.05968899]
[ 3.07686272]
[-1.19205355]
[-1.54193042]
[-4.46625195]
[-4.83257391]
[ 0.41097987]
[-4.00972237]
[ 2.56088925]
[ 3.2564741 ]]
```

1.1 A)

Estimate the value of the regression coefficients by using the analytical solution for the least squares estimation problem.

```

In [73]: time_start = time.clock()
         beta_ls_exact=np.dot(np.dot(np.linalg.inv(np.dot(np.transpose(X),X)),np.tr
         time_elapsed = (time.clock() - time_start)
         print('time elapsed=',time_elapsed)
         print(beta_ls_exact)
```

```
('time elapsed=', 0.0023641579900868237)
[[ 1.73944484]
 [ 2.11849576]
 [-3.3627126 ]
 [ 2.49636401]
 [-2.84054722]
 [ 2.02156168]
 [-2.86232325]
 [ 0.35178602]
 [ 1.543078  ]
 [ 3.75664986]
 [-3.10144696]
 [-0.49563234]
 [ 4.07457553]
 [ 3.19734595]
 [ 4.01312126]
 [ 1.81554646]
 [ 2.73727512]
 [ 4.86453595]
 [ 4.78753664]
 [-2.28646111]
 [-2.60981044]
 [-4.13475238]
 [-2.42150444]
 [ 4.51678882]
 [-4.76576907]
 [ 4.55297872]
 [ 4.76157291]
 [-0.63205529]
 [-4.11615598]
 [ 1.24171108]
 [-3.95975354]
 [-1.08472859]
 [-0.25454118]
 [-2.45531301]
 [ 2.04235535]
 [ 4.90231666]
 [ 2.83335418]
 [ 0.28134911]
 [ 4.84884805]
 [ 0.42513108]
 [-0.79616354]
 [ 3.25057585]
 [-1.57517965]
 [ 0.0961234  ]
 [ 2.17431404]
 [ 1.82505696]
 [ 4.22074613]
```

[2.44749707]
[-0.70609706]
[-1.94542769]
[4.95931038]
[2.40542927]
[-2.24506929]
[4.64658585]
[-2.99305659]
[-2.30420422]
[-4.70322873]
[2.55568737]
[3.39306073]
[4.69860549]
[4.21740776]
[-3.46522597]
[-4.70731893]
[-0.38375324]
[0.15328952]
[4.26222374]
[4.69090512]
[1.21393678]
[-4.91855814]
[0.60281507]
[3.99072754]
[2.98359187]
[4.1058776]
[4.92828094]
[-2.10718704]
[3.9875125]
[1.06448384]
[-1.58946094]
[2.42858799]
[-4.70019775]
[-2.99818303]
[2.16842197]
[4.62660575]
[-3.42353357]
[3.24789051]
[-4.59425366]
[-0.09865654]
[4.73365641]
[-4.82251921]
[2.85609984]
[-2.43892912]
[0.05248573]
[3.066192]
[-1.18456813]
[-1.53919184]

```

[-4.4824774 ]
[-4.83151957]
[ 0.41117392]
[-4.02762808]
[ 2.57962647]
[ 3.24912254]]

```

1.2 B

```

In [76]: import matplotlib.pyplot as plt
          %matplotlib inline
          lambda_est=[]
          sigma=np.arange(1,10, 0.5)

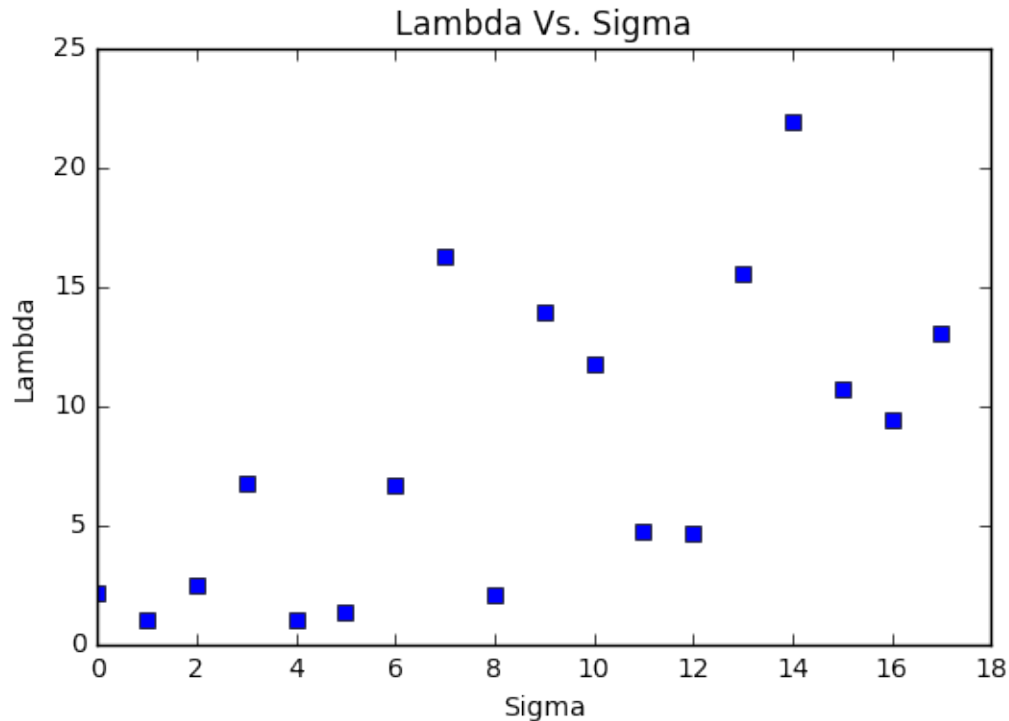
          for i in sigma :
              error=np.random.normal(0, i, (nsample,1)) # Normal random error with o
              Y=np.dot(X, beta)+error
              beta_ls_exact=np.dot(np.dot(np.linalg.inv(np.dot(np.transpose(X),X)),n
              lambda_est.append((np.linalg.norm(beta-beta_ls_exact,ord=2)/np.linalg.

          print lambda_est

          plt.plot(lambda_est, 'bs')
          plt.xlabel('Sigma')
          plt.ylabel('Lambda')
          plt.title('Lambda Vs. Sigma')
          plt.show()

[2.0977724073092068, 1.0434809952479307, 2.4665810510725774, 6.714342056905366, 0.9

```



1.3 C)

Estimate the value of the regression coefficients (least squares) by using the tool minimize from the python package Scipy.optimize. Try at least three available solvers and compare their performance (iterations, function, gradient and hessian evaluations as well as total computational time).

```
In [77]: # c : Newton-CG
         from scipy.optimize import minimize

         def least_sq_reg_der(beta_ls,X,Y):
             beta_ls=np.matrix(beta_ls)
             pp=-2*np.dot(np.transpose(Y-np.dot(X,np.transpose(beta_ls))),X)
             return np.squeeze(np.asarray(pp))

         def least_sq_reg_hess(beta_ls,X,Y):
             return 2*np.dot(np.transpose(X),X)

         time_start = time.clock()
         res = minimize(least_sq_reg, beta_ls0, args=(X,Y), method = 'Newton-CG', jac=least_sq_reg_der)
         time_elapsed = (time.clock() - time_start)
```

```

print('time elapsed=',time_elapsed)
print(res.x)
np.linalg.inv(least_sq_reg_hess(res.x,X,Y))
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-res.x,ord=2)/np.

```

Optimization terminated successfully.

Current function value: 84919.433779

Iterations: 13

Function evaluations: 15

Gradient evaluations: 27

Hessian evaluations: 13

('time elapsed=', 0.01405822801461909)

```

[-2.82008067  2.18757466 -3.33272327  2.40049448 -2.91963727  2.00327112
 -2.92131844  0.38737346  1.41670753  3.80403476 -3.04939229 -0.58693961
  4.13023388  3.2712188  4.06065034  1.86865285  2.90094534  5.10282956
  4.75282604 -2.3773995 -2.74905144 -3.94538837 -2.45767374  4.49870427
 -4.74963116  4.5581906  4.65647517 -0.59536712 -4.19275214  1.24511272
 -3.89736906 -1.04011949 -0.2648055 -2.28654196  2.16801341  4.92252474
  2.94042701  0.53332892  4.79756196  0.43628567 -0.86060568  3.35301637
 -1.41157385  0.06995832  2.14622713  1.82173066  4.29591559  2.45075603
 -0.63613901 -1.91125291  4.94260168  2.30714525 -2.39602297  4.6346139
 -3.05933576 -2.51372002 -4.71725957  2.5502426  3.42703598  4.78687509
  4.21202156 -3.54248362 -4.50429966 -0.48099936  0.13159525  4.24485933
  4.77287047  1.38593122 -4.99300658  0.74922483  3.95908503  2.79126195
  3.99975694  4.84343316 -2.07182145  3.98490325  1.19556984 -1.65011529
  2.41928262 -4.77430334 -2.80623367  2.12898233  4.6608253 -3.45711524
  3.31431846 -4.59174864 -0.0981681  4.70477634 -4.95014163  2.87792685
 -2.4599172  0.12189708  2.88425646 -1.36533708 -1.58816479 -4.39826269
 -4.80547961  0.56007547 -3.78563627  2.50914512  3.30404331]
('error=', 1.3659231966901152e-06)

```

In [78]: # c: dogleg

```

from scipy.optimize import minimize

```

```

def least_sq_reg_der(beta_ls,X,Y):
    beta_ls=np.matrix(beta_ls)
    pp=-2*np.dot(np.transpose(Y-np.dot(X,np.transpose(beta_ls))),X)
    return np.squeeze(np.asarray(pp))

```

```

def least_sq_reg_hess(beta_ls,X,Y):
    return 2*np.dot(np.transpose(X),X)

```

```

time_start = time.clock()

```



```

res = minimize(least_sq_reg, beta_ls0, args=(X,Y), method = 'dogleg', jac=J)
time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print(res.x)
np.linalg.inv(least_sq_reg_hess(res.x,X,Y))
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-res.x,ord=2)/np.

```

Optimization terminated successfully.

Current function value: 84919.433768

Iterations: 6

Function evaluations: 7

Gradient evaluations: 7

Hessian evaluations: 6

('time elapsed=', 0.008869112425600179)

```

[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
 -2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
  4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
  4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
 -4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
 -3.8973661 -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
  2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
 -1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
 -0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
 -3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769
  4.21201842 -3.5424867 -4.50429917 -0.48099319  0.1316062  4.24485554
  4.7728776  1.38593473 -4.99300561  0.74922867  3.95908786  2.79126176
  3.99975326  4.84343426 -2.07181554  3.98489739  1.19557716 -1.65011212
  2.41928082 -4.7743059 -2.80624142  2.12898243  4.66082827 -3.45711962
  3.31431908 -4.59174474 -0.0981666  4.70477382 -4.95014286  2.87792284
 -2.45991397  0.12189604  2.88425693 -1.36533547 -1.58816079 -4.39826434
 -4.8054921  0.56007872 -3.78564008  2.50915003  3.30404243]

```

('error=', 4.2990732639375794e-12)

In [79]: `from scipy.optimize import minimize`

```

def least_sq_reg(beta_ls, X, Y):
    beta_ls=np.matrix(beta_ls)
    z=Y-X*np.transpose(beta_ls)
    return np.dot(np.transpose(z),z) #sum of the square errors, value of t

```

In [80]: `# c :BFGS`

```

beta_ls0 = np.zeros(nvariables+1) #we need to initiate our initial point
time_start = time.clock()
res = minimize(least_sq_reg, beta_ls0, args=(X,Y), method = 'BFGS', option
time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)

```

```

print (res.x)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-res.x,ord=2)/np.

Warning: Desired error not necessarily achieved due to precision loss.
Current function value: 84919.433768
Iterations: 131
Function evaluations: 24515
Gradient evaluations: 238
('time elapsed=', 1.4608064334461233)
[-2.8200985  2.18757251 -3.33272404  2.40050027 -2.9196443  2.00327308
 -2.92131936  0.3873754  1.41670726  3.80403212 -3.04939138 -0.5869397
  4.13024253  3.27121944  4.0606465  1.86865365  2.90094716  5.10282556
  4.75282723 -2.3774021 -2.74905506 -3.94539385 -2.45767775  4.49870628
 -4.74962365  4.5581867  4.65647564 -0.59536055 -4.19274501  1.2451148
 -3.89736611 -1.04011893 -0.2648137 -2.28654166  2.1680069  4.92252054
  2.94043061  0.53332426  4.79756441  0.43627895 -0.86060973  3.35301358
 -1.411572  0.06996955  2.14622612  1.82173174  4.29591722  2.45076115
 -0.63614186 -1.91125438  4.94259787  2.30714665 -2.39602388  4.63460324
 -3.05933111 -2.51371897 -4.71726272  2.55024273  3.42703609  4.78687685
  4.2120184 -3.54248668 -4.50429913 -0.48099318  0.13160623  4.24485561
  4.77287762  1.38593473 -4.99300558  0.7492287  3.95908792  2.79126178
  3.99975324  4.84343427 -2.07181555  3.98489735  1.19557716 -1.65011214
  2.41928088 -4.77430592 -2.80624141  2.12898241  4.66082837 -3.45711963
  3.31431909 -4.59174468 -0.09816652  4.70477388 -4.95014286  2.8779229
 -2.45991397  0.12189603  2.8842569 -1.36533542 -1.58816079 -4.39826433
 -4.80549213  0.56007879 -3.78564002  2.50914997  3.30404246]
('error=', 2.1148376814635847e-07)

```

In [81]: # c :SLSQP

```

beta_ls0 = np.zeros(nvariables+1) #we need to initiate our initial point
time_start = time.clock()
res = minimize(least_sq_reg, beta_ls0, args=(X,Y), method = 'SLSQP', options={})
time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print (res.x)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-res.x,ord=2)/np.

Optimization terminated successfully.      (Exit mode 0)
Current function value: [[ 84919.43438612]]
Iterations: 62
Function evaluations: 6628
Gradient evaluations: 62
('time elapsed=', 0.38456415182736237)
[-2.81999282  2.18757499 -3.33274337  2.40049986 -2.91964145  2.00328562
 -2.92135673  0.38740984  1.41667548  3.8040489 -3.04938633 -0.58696508
  4.13026249  3.2712273  4.06064785  1.86863405  2.90094692  5.10280418

```

```

4.75279975 -2.37740227 -2.74908372 -3.94539617 -2.45771762 4.49873764
-4.7495816 4.55816327 4.65648354 -0.5953383 -4.19275952 1.24510698
-3.89736525 -1.04011663 -0.26475866 -2.28654486 2.16800036 4.92250585
2.94045632 0.53331927 4.79756062 0.43632093 -0.86062636 3.35304649
-1.41155463 0.07002751 2.14629284 1.82174526 4.29594552 2.45075327
-0.63611914 -1.9112522 4.94254981 2.3071427 -2.39602753 4.63456411
-3.05933373 -2.51372557 -4.71728558 2.55019751 3.42705259 4.78684688
4.21203778 -3.54249077 -4.50430907 -0.48098153 0.13159908 4.24482322
4.77287968 1.38597311 -4.99303748 0.74922122 3.95905633 2.79125096
3.99978166 4.84342029 -2.07179286 3.98486089 1.19553417 -1.6501765
2.41921376 -4.77430548 -2.80619574 2.12901886 4.66084734 -3.45709624
3.31433304 -4.59175443 -0.09815162 4.70480099 -4.95015551 2.87791769
-2.45999135 0.12189591 2.88425535 -1.36529894 -1.58814319 -4.39827985
-4.80542618 0.56006648 -3.78563869 2.50912917 3.30408075]
('error=', 9.0791862297303001e-06)

```

In [82]: # c : CG

```

beta_ls0 = np.zeros(nvariables+1) #we need to initiate our initial point
time_start = time.clock()
res = minimize(least_sq_reg, beta_ls0, args=(X,Y), method = 'CG', options
time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print (res.x)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-res.x,ord=2)/np.

```

Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 84919.438062

Iterations: 243

Function evaluations: 55632

Gradient evaluations: 540

('time elapsed=', 2.9277961149491603)

```

[-2.85789181 2.1876595 -3.33265676 2.40058637 -2.91958424 2.0033621
-2.92124797 0.38743794 1.41677483 3.8040492 -3.04932774 -0.58686475
4.13034603 3.27131329 4.06073374 1.86871524 2.90101457 5.10290849
4.75291173 -2.37733101 -2.74897404 -3.94534963 -2.45755214 4.4987615
-4.74954624 4.55825218 4.65654622 -0.59529921 -4.19269757 1.24517746
-3.89728324 -1.04006126 -0.26472422 -2.28641039 2.16808449 4.92261235
2.94051853 0.5333913 4.7976289 0.43632826 -0.86052493 3.35309622
-1.41149609 0.07004405 2.14633029 1.8218001 4.2959728 2.45083466
-0.63607714 -1.91116057 4.94266668 2.30720999 -2.39592101 4.63465029
-3.05928473 -2.5136702 -4.7171507 2.5503713 3.42710086 4.78695988
4.21207411 -3.54236985 -4.50420913 -0.48090851 0.13170382 4.24491855
4.77296519 1.3860117 -4.99294112 0.74922996 3.95915735 2.79134114
3.99980412 4.84347789 -2.0717737 3.98497698 1.19565988 -1.65003571
2.41931526 -4.77427489 -2.80616743 2.12906819 4.66091626 -3.45704269
3.3143886 -4.59171885 -0.09806332 4.70488323 -4.95006486 2.87796457

```

```
-2.45985323  0.12197447  2.88435942 -1.36521473 -1.58804764 -4.39818139
-4.8054392   0.56020827 -3.78554284  2.50927804  3.30413515]
('error=', 0.0011659627184107803)
```

```
In [4]: import sympy
        from pandas import DataFrame
        from sympy import *

        # For Final report we need to put 'Current function value', 'Iteration', 'Fu
        ### Fill value instead of 1 2 3 .....
        t = [[84919.433779 ,13 ,15 ,27 ,13 ,0.014 ,1.37e-6 ], #Newton-CG
              [84919.434768 , 6, 7 ,7 ,6 , 0.0089 ,4.30e-12 ], #dogleg
              [84919.433758 , 131,24515 ,238 ,"NA" ,1.46 ,2.11e-7 ], #BFGS
              [84919.434386 ,62 ,6628 ,62 ,"NA" ,0.38 ,9.08e-6 ], #SLSQP
              [84919.438062 ,243 ,55632 ,540 ,"NA" ,2.93 ,0.0012 ]] #CG

        table = DataFrame(t, index=['Newton-CG', 'dogleg', 'BFGS', 'SLSQP', 'CG'],

        table
```

```
Out[4]:
```

	Current function value	Iteration	Function evaluations	\
Newton-CG	84919.433779	13	15	
dogleg	84919.434768	6	7	
BFGS	84919.433758	131	24515	
SLSQP	84919.434386	62	6628	
CG	84919.438062	243	55632	

	Gradient evaluations	Hessian evaluations	Time elapsed	\
Newton-CG	27	13	0.0140	
dogleg	7	6	0.0089	
BFGS	238	NA	1.4600	
SLSQP	62	NA	0.3800	
CG	540	NA	2.9300	

	Error
Newton-CG	1.370000e-06
dogleg	4.300000e-12
BFGS	2.110000e-07
SLSQP	9.080000e-06
CG	1.200000e-03

1.4 D)

Considering again the least squares estimation problem, estimate the value of the regression coefficients by implementing the:

- i. Gradient method
- ii. Newton method
- iii. Quasi-Newton method

Consider a like search technique to improve the algorithm convergence, e.x., Armijo rule. Compare the performance of these algorithms (iterations, function, gradient and hessian evaluations as well as total computational time).

```
In [84]: #definition of OF #objective function
def least_sq_reg(beta_ls, X, Y):
    beta_ls=np.matrix(beta_ls)
    z=Y-X*np.transpose(beta_ls)
    return np.transpose(z)*z

#definition of Gradient
def least_sq_reg_der(beta_ls,X,Y):
    beta_ls=np.matrix(beta_ls)
    pp=-2*np.transpose(Y-X*np.transpose(beta_ls))*X
    aa= np.squeeze(np.asarray(pp))
    return aa

#definition of hessian
def least_sq_reg_hess(beta_ls,X,Y):
    ss=2*np.dot(np.transpose(X),X)
    return ss
```

1.5 i) Gradient method

- From an initial iterate x_0
 - Compute search (descent) directions p_k
 - Far from the solution, compute a steplength $\alpha_k > 0$
 - Movement:

$$x_{k+1} = x_k + \alpha_k p_k$$

Until convergence to a local solution

```
In [97]: ### With Armijo Rule

(a,b)=X.shape
beta_lsg=np.zeros(b) #initial value for beta
#alpha= 0.0000001
n_iter=100000 #maximim nunber iteration
OF_iter=np.zeros(n_iter)
tol_iter=np.zeros(n_iter)
alpha_iter=np.zeros(n_iter)
i=0;
tol=1000; # Tolerance
epsilon=1e-3;
```

```

time_start = time.clock()
while (i <= n_iter-2) and (tol>epsilon):
    i=i+1
    grad=least_sq_reg_der(beta_lsg,X,Y) #this function gives us the value
    ddirect=-grad
    #####
    # Armijo Rule
    sigma=0.1
    beta=0.5
    alpha=1
    while (least_sq_reg(beta_lsg+alpha*ddirect,X,Y)> least_sq_reg(beta_lsg)):
        alpha=alpha*beta
        #####
        beta_lsg=beta_lsg+alpha*ddirect
        OF_iter[i]=least_sq_reg(beta_lsg,X,Y) # Objective Function ---Residual
        tol=np.linalg.norm(grad,ord=2)
        tol_iter[i]=tol

time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print('iterations',i)
print(OF_iter[i])
print(beta_lsg)
print(np.transpose(beta_ls_exact))
print('Tolerance=',tol)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-beta_lsg,ord=2)/

('time elapsed=', 215.57388964362326)
('iterations', 99999)
84929.3073365
[-1.00552324  2.18336177 -3.33600171  2.39643687 -2.92246059  1.99875166
 -2.92468279  0.38429239  1.41333943  3.80317979 -3.05248791 -0.59071866
  4.12524893  3.26656723  4.05649917  1.86564815  2.8978303  5.09889964
  4.74850713 -2.38069513 -2.75260615 -3.94742467 -2.46350012  4.49601821
 -4.75297247  4.55537918  4.65282515 -0.59800657 -4.19509426  1.24210405
 -3.90135134 -1.04271539 -0.26901538 -2.29302437  2.16423839  4.91800017
  2.93592439  0.53012202  4.79433681  0.43390435 -0.86465576  3.34852329
 -1.41526794  0.06632625  2.14111084  1.81840842  4.29306879  2.44703809
 -0.63913949 -1.91576422  4.93924581  2.30399171 -2.40088197  4.63240874
 -3.06169684 -2.5159778  -4.72248466  2.54384261  3.42392076  4.7828991
  4.20934812 -3.54783455 -4.50863292 -0.48497834  0.1268127  4.24195353
  4.76853478  1.38218161 -4.99625161  0.74897886  3.95538557  2.78762389
  3.99727475  4.84105212 -2.0735365  3.98123654  1.19165168 -1.65348411
  2.41727442 -4.77572941 -2.80935509  2.12472609  4.65639364 -3.46088663
  3.31067667 -4.59263375 -0.10338909  4.6996278  -4.95383034  2.87600641
 -2.46276516  0.11816647  2.87931377 -1.37122449 -1.59362318 -4.40189991
 -4.80763108  0.55374605 -3.79003875  2.50319568  3.29974066]

```

```

[[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
 -2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
  4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
  4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
 -4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
 -3.8973661  -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
  2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
 -1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
 -0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
 -3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769
  4.21201842 -3.5424867  -4.50429917 -0.48099319  0.1316062  4.24485554
  4.7728776  1.38593473 -4.99300561  0.74922867  3.95908786  2.79126176
  3.99975326  4.84343426 -2.07181554  3.98489739  1.19557716 -1.65011212
  2.41928082 -4.7743059  -2.80624142  2.12898243  4.66082827 -3.45711962
  3.31431908 -4.59174474 -0.0981666  4.70477382 -4.95014286  2.87792284
 -2.45991397  0.12189604  2.88425693 -1.36533547 -1.58816079 -4.39826434
 -4.8054921  0.56007872 -3.78564008  2.50915003  3.30404243]]
('Tolerance=', 11.957228882608536)
('error=', 0.056169323979137355)

```

In [98]: *### Without Armijo Rule*

```

(a,b)=X.shape
beta_lsg=np.zeros(b) #initial value for beta
alpha= 0.0000001
n_iter=100000 #maximim nunber iteration
OF_iter=np.zeros(n_iter)
tol_iter=np.zeros(n_iter)
alpha_iter=np.zeros(n_iter)
i=0;
tol=1000; # Tolerance
epsilon=1e-3;

time_start = time.clock()
while (i <= n_iter-2) and (tol>epsilon):
    i=i+1
    grad=least_sq_reg_der(beta_lsg,X,Y) #this function gives us the value
    ddirect=-grad
    beta_lsg=beta_lsg+alpha*ddirect
    OF_iter[i]=least_sq_reg(beta_lsg,X,Y) # Objective Function ---Residual
    tol=np.linalg.norm(grad,ord=2)
    tol_iter[i]=tol

time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print('iterations',i)
print(OF_iter[i])

```

```

print(beta_lsg)
print(np.transpose(beta_ls_exact))
print('Tolerance=',tol)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-beta_lsg,ord=2),

('time elapsed=', 14.0880259515543)
('iterations', 99999)
84942.527301
[-0.0449488  2.18113141 -3.33773806  2.39428447 -2.92395271  1.99635687
 -2.92646459  0.38265903  1.41155527  3.80272724 -3.05412837 -0.59272041
  4.12260421  3.26410322  4.05430245  1.86405585  2.89617906  5.09682015
  4.74621892 -2.38243967 -2.75448725 -3.94850098 -2.46658356  4.49459399
 -4.75474652  4.55389168  4.65089142 -0.59940854 -4.19633916  1.24050891
 -3.90346229 -1.04409118 -0.27124091 -2.29645735  2.16224217  4.91560592
  2.93353771  0.52842552  4.79262692  0.43264598 -0.86679887  3.34614498
 -1.41722575  0.06439638  2.13840169  1.8166479  4.29155963  2.44506596
 -0.64072767 -1.91815287  4.93747005  2.30232025 -2.40345492  4.63124573
 -3.06295048 -2.51717478 -4.72525025  2.54045333  3.4222703  4.78079204
  4.20793326 -3.55066681 -4.51092836 -0.48708925  0.12427388  4.240416
  4.76623454  1.38019354 -4.99797124  0.74884532  3.95342441  2.78569681
  3.99596141  4.83978978 -2.07444881  3.9792973  1.18957239 -1.65527045
  2.41621101 -4.77648427 -2.81100466  2.1224716  4.65404478 -3.46288209
  3.30874722 -4.59310568 -0.106155  4.69690238 -4.95578365  2.87499064
 -2.46427576  0.11619089  2.87669574 -1.37434324 -1.59651609 -4.40382578
 -4.80876468  0.55039243 -3.7923686  2.50004233  3.29746211]
[[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
 -2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
  4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
  4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
 -4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
 -3.8973661 -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
  2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
 -1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
 -0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
 -3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769
  4.21201842 -3.5424867 -4.50429917 -0.48099319  0.1316062  4.24485554
  4.7728776  1.38593473 -4.99300561  0.74922867  3.95908786  2.79126176
  3.99975326  4.84343426 -2.07181554  3.98489739  1.19557716 -1.65011212
  2.41928082 -4.7743059 -2.80624142  2.12898243  4.66082827 -3.45711962
  3.31431908 -4.59174474 -0.0981666  4.70477382 -4.95014286  2.87792284
 -2.45991397  0.12189604  2.88425693 -1.36533547 -1.58816079 -4.39826434
 -4.8054921  0.56007872 -3.78564008  2.50915003  3.30404243]]
('Tolerance=', 16.639447266111439)
('error=', 0.085955155074071907)

```

In []:

1.6 ii) Newtons method

→ From an initial iterate x_0

→ Compute search (descent) directions $p_k = -(\nabla^2 f(x_k))^{-1} \nabla f(x_k)$, whenever $\nabla^2 f(x_k)$ is non-singular.

→ Far from the solution, compute a steplength $\alpha_k > 0$

→ Movement:

$$x_{k+1} = x_k + \alpha_k p_k$$

Until convergence to a local solution

In [99]: *### With Armijo Rule*

```
(a,b)=X.shape
beta_lsg=np.zeros(b) #initial value for beta
#alpha= 0.0000001
n_iter=100000 #maximim nunber iteration
OF_iter=np.zeros(n_iter)
tol_iter=np.zeros(n_iter)
alpha_iter=np.zeros(n_iter)
i=0;
tol=1000; # Tolerance
epsilon=1e-3;

time_start = time.clock()
while (i <= n_iter-2) and (tol>epsilon):
    i=i+1
    grad=least_sq_reg_der(beta_lsg,X,Y)
    hess=least_sq_reg_hess(beta_lsg,X,Y) #this function gives us the value
    ddirect=-(np.dot(np.linalg.inv(hess), grad))
    #####
    # Armijo Rule
    sigma=0.1
    beta=0.5
    alpha=1
    while (least_sq_reg(beta_lsg+alpha*ddirect,X,Y)> least_sq_reg(beta_lsg)):
        alpha=alpha*beta
    #####
    beta_lsg=beta_lsg+alpha*ddirect
    OF_iter[i]=least_sq_reg(beta_lsg,X,Y)
    tol=np.linalg.norm(hess,ord=2)
    tol_iter[i]=tol

time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print('iterations',i)
print(OF_iter[i])
print(beta_lsg)
print(np.transpose(beta_ls_exact))
```

```

print('Tolerance=',tol)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-beta_lsg,ord=2)/

('time elapsed=', 293.5455931344186)
('iterations', 99999)
84919.4337684
[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
 -2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
  4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
  4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
 -4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
 -3.8973661  -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
  2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
 -1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
 -0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
 -3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769
  4.21201842 -3.5424867  -4.50429917 -0.48099319  0.1316062  4.24485554
  4.7728776  1.38593473 -4.99300561  0.74922867  3.95908786  2.79126176
  3.99975326  4.84343426 -2.07181554  3.98489739  1.19557716 -1.65011212
  2.41928082 -4.7743059  -2.80624142  2.12898243  4.66082827 -3.45711962
  3.31431908 -4.59174474 -0.0981666  4.70477382 -4.95014286  2.87792284
 -2.45991397  0.12189604  2.88425693 -1.36533547 -1.58816079 -4.39826434
 -4.8054921  0.56007872 -3.78564008  2.50915003  3.30404243]
[[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
 -2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
  4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
  4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
 -4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
 -3.8973661  -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
  2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
 -1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
 -0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
 -3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769
  4.21201842 -3.5424867  -4.50429917 -0.48099319  0.1316062  4.24485554
  4.7728776  1.38593473 -4.99300561  0.74922867  3.95908786  2.79126176
  3.99975326  4.84343426 -2.07181554  3.98489739  1.19557716 -1.65011212
  2.41928082 -4.7743059  -2.80624142  2.12898243  4.66082827 -3.45711962
  3.31431908 -4.59174474 -0.0981666  4.70477382 -4.95014286  2.87792284
 -2.45991397  0.12189604  2.88425693 -1.36533547 -1.58816079 -4.39826434
 -4.8054921  0.56007872 -3.78564008  2.50915003  3.30404243]]
('Tolerance=', 4972171.1853834139)
('error=', 4.2435594497630421e-12)

```

In [100]: *### Without Armijo Rule*

```

(a,b)=X.shape
beta_lsg=np.zeros(b) #initial value for beta

```

```

alpha= 0.0000001
n_iter=100000 #maximim nunber iteration
OF_iter=np.zeros(n_iter)
tol_iter=np.zeros(n_iter)
alpha_iter=np.zeros(n_iter)
i=0;
tol=1000; # Tolerance
epsilon=1e-3;

time_start = time.clock()
while (i <= n_iter-2) and (tol>epsilon):
    i=i+1
    grad=least_sq_reg_der(beta_lsg,X,Y)
    hess=least_sq_reg_hess(beta_lsg,X,Y) #this function gives us the value of hessian
    ddirect=-(np.dot(np.linalg.inv(hess), grad))
    beta_lsg=beta_lsg+alpha*ddirect
    OF_iter[i]=least_sq_reg(beta_lsg,X,Y)
    tol=np.linalg.norm(hess,ord=2)
    tol_iter[i]=tol

time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print('iterations',i)
print(OF_iter[i])
print(beta_lsg)
print(np.transpose(beta_ls_exact))
print('Tolerance=',tol)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-beta_lsg,ord=2))

('time elapsed=', 192.15463624136464)
('iterations', 99999)
84611989.3714
[-0.0280601  0.0217665 -0.03316083  0.02388514 -0.02905066  0.0199327
-0.02906733  0.00385441  0.01409633  0.03785038 -0.03034165 -0.00584009
 0.04109619  0.03254885  0.04040371  0.01859323  0.02886462  0.05077346
 0.04729095 -0.02365531 -0.02735328 -0.03925694 -0.02445406  0.04476243
-0.04725908  0.04535427  0.04633225 -0.00592388 -0.0417181  0.01238898
-0.03877906 -0.01034925 -0.00263491 -0.02275125  0.02157182  0.04897941
 0.02925748  0.00530661  0.04773609  0.00434101 -0.00856313  0.03336271
-0.01404524  0.0006962  0.0213551  0.01812635  0.04274467  0.02438524
-0.00632965 -0.01901711  0.04917918  0.02295627 -0.0238406  0.04611462
-0.03044055 -0.02501167 -0.04693708  0.02537509  0.03409924  0.04762975
 0.04190987 -0.03524798 -0.04481808 -0.00478591  0.00130949  0.0422366
 0.04749046  0.01379014 -0.04968074  0.00745488  0.03939319  0.02777324
 0.03979782  0.0481925  -0.02061471  0.03965  0.01189607 -0.01641873
 0.02407201 -0.04750467 -0.02792229  0.02118352  0.04637556 -0.03439857
 0.0329777  -0.04568817 -0.00097676  0.04681282 -0.04925426  0.02863553
-0.02447631  0.00121287  0.02869855 -0.01358518 -0.01580231 -0.04376303]

```

```

-0.04781497  0.00557282 -0.03766738  0.02496621  0.03287545]
[[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
 -2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
  4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
  4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
 -4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
 -3.8973661  -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
  2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
 -1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
 -0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
 -3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769
  4.21201842 -3.5424867  -4.50429917 -0.48099319  0.1316062  4.24485554
  4.7728776  1.38593473 -4.99300561  0.74922867  3.95908786  2.79126176
  3.99975326  4.84343426 -2.07181554  3.98489739  1.19557716 -1.65011212
  2.41928082 -4.7743059  -2.80624142  2.12898243  4.66082827 -3.45711962
  3.31431908 -4.59174474 -0.0981666  4.70477382 -4.95014286  2.87792284
 -2.45991397  0.12189604  2.88425693 -1.36533547 -1.58816079 -4.39826434
 -4.8054921  0.56007872 -3.78564008  2.50915003  3.30404243]]
('Tolerance=', 4972171.1853834139)
('error=', 99.501828333603413)

```

1.7 iii) Quasi-Newton method

In [101]: *### With Armijo Rule*

```

(a,b)=X.shape
beta_quasi=np.zeros(b) #initial value for beta
#alpha= 0.0000001
n_iter=100000 #maximim nunber iteration
OF_iter=np.zeros(n_iter)
tol_iter=np.zeros(n_iter)
alpha_iter=np.zeros(n_iter)
i=0;
tol=1000; # Tolerance
epsilon=1e-5;

time_start = time.clock()
while (i <= n_iter-2) and (tol>epsilon):
    i=i+1
    grad=least_sq_reg_der(beta_quasi,X,Y)
    if (i==1):
        grad=least_sq_reg_der(beta_quasi,X,Y)
        B=least_sq_reg_hess(beta_quasi,X,Y)
    else:

```

```

grad_previous=grad
grad=least_sq_reg_der(beta_quasi,X,Y)
y= grad - grad_previous
s= beta_quasi - beta_quasi_previous

a=np.dot(B, s)
b=np.transpose(np.dot(B, s))
c=np.dot(np.dot(np.transpose(s), B), s)
d=np.dot(y, np.transpose(y))
e=np.dot(np.transpose(y), s)

### Page 44 Formulas
B= B + np.dot(y-a, np.transpose(y-a))/np.dot(np.transpose(y-a), s)
# B= B - (np.dot(a,b)/c)+(d/e) (BFGS Formula)
ddirect=-(np.dot(np.linalg.inv(B), grad))

#####
# Armijo Rule
sigma=0.1
beta=0.5
alpha=1
while (least_sq_reg(beta_quasi+alpha*ddirect,X,Y)> least_sq_reg(beta_
alpha=alpha*beta
#####
beta_quasi_previous= beta_quasi
beta_quasi=beta_quasi_previous + alpha*ddirect
OF_iter[i]=least_sq_reg(beta_quasi,X,Y)
tol=np.linalg.norm(grad,ord=2)
tol_iter[i]=tol

time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print('iterations',i)
print(OF_iter[i])
print(beta_quasi)
print(np.transpose(beta_ls_exact))
print('Tolerance=',tol)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-beta_quasi,ord=2))

('time elapsed=', 0.0035460236395010725)
('iterations', 2)
84919.4337684
[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
 -2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
  4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
  4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
 -4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
 -3.8973661 -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054

```

```

2.94043056 0.5333243 4.79756441 0.43627901 -0.86060974 3.35301358
-1.41157203 0.06996948 2.14622611 1.8217317 4.29591723 2.45076112
-0.63614187 -1.91125437 4.94259785 2.30714669 -2.39602395 4.63460325
-3.05933115 -2.51371899 -4.71726272 2.55024272 3.42703614 4.7868769
4.21201842 -3.5424867 -4.50429917 -0.48099319 0.1316062 4.24485554
4.7728776 1.38593473 -4.99300561 0.74922867 3.95908786 2.79126176
3.99975326 4.84343426 -2.07181554 3.98489739 1.19557716 -1.65011212
2.41928082 -4.7743059 -2.80624142 2.12898243 4.66082827 -3.45711962
3.31431908 -4.59174474 -0.0981666 4.70477382 -4.95014286 2.87792284
-2.45991397 0.12189604 2.88425693 -1.36533547 -1.58816079 -4.39826434
-4.8054921 0.56007872 -3.78564008 2.50915003 3.30404243]
[[-2.82009166 2.18757262 -3.33272408 2.40050035 -2.91964431 2.00327309
-2.92131935 0.38737545 1.41670728 3.80403217 -3.04939149 -0.58693972
4.13024247 3.27121943 4.06064646 1.86865365 2.9009471 5.1028255
4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779 4.49870622
-4.74962366 4.55818668 4.65647563 -0.59536061 -4.19274504 1.24511478
-3.8973661 -1.04011885 -0.26481368 -2.28654173 2.16800687 4.92252054
2.94043056 0.5333243 4.79756441 0.43627901 -0.86060974 3.35301358
-1.41157203 0.06996948 2.14622611 1.8217317 4.29591723 2.45076112
-0.63614187 -1.91125437 4.94259785 2.30714669 -2.39602395 4.63460325
-3.05933115 -2.51371899 -4.71726272 2.55024272 3.42703614 4.7868769
4.21201842 -3.5424867 -4.50429917 -0.48099319 0.1316062 4.24485554
4.7728776 1.38593473 -4.99300561 0.74922867 3.95908786 2.79126176
3.99975326 4.84343426 -2.07181554 3.98489739 1.19557716 -1.65011212
2.41928082 -4.7743059 -2.80624142 2.12898243 4.66082827 -3.45711962
3.31431908 -4.59174474 -0.0981666 4.70477382 -4.95014286 2.87792284
-2.45991397 0.12189604 2.88425693 -1.36533547 -1.58816079 -4.39826434
-4.8054921 0.56007872 -3.78564008 2.50915003 3.30404243]]
('Tolerance=', 2.987721920840585e-06)
('error=', 6.429578978306502e-13)

```

In [102]: *### Without Armijo Rule*

```

(a,b)=X.shape
beta_quasi=np.zeros(b) #initial value for beta
alpha= 0.0000001
n_iter=100000 #maximim nunber iteration
OF_iter=np.zeros(n_iter)
tol_iter=np.zeros(n_iter)
alpha_iter=np.zeros(n_iter)
i=0;
tol=1000; # Tolerance
epsilon=1e-5;

```

```

time_start = time.clock()
while (i <= n_iter-2) and (tol>epsilon):
    i=i+1
    grad=least_sq_reg_der(beta_quasi,X,Y)
    if (i==1):
        grad=least_sq_reg_der(beta_quasi,X,Y)
        B=least_sq_reg_hess(beta_quasi,X,Y)
    else:
        grad_previous=grad
        grad=least_sq_reg_der(beta_quasi,X,Y)
        y= grad - grad_previous
        s= beta_quasi - beta_quasi_previous

        a=np.dot(B, s)
        b=np.transpose(np.dot(B, s))
        c=np.dot(np.dot(np.transpose(s), B), s)
        d=np.dot(y, np.transpose(y))
        e=np.dot(np.transpose(y), s)

        ### Page 44 Formulas

        B= B + np.dot(y-a, np.transpose(y-a))/np.dot(np.transpose(y-a), s)

        # B= B - (np.dot(a,b)/c)+(d/e)    (BFGS Formula)
        ddirect=-(np.dot(np.linalg.inv(B), grad))

        beta_quasi_previous= beta_quasi
        beta_quasi=beta_quasi_previous + alpha*ddirect
        OF_iter[i]=least_sq_reg(beta_quasi,X,Y)
        tol=np.linalg.norm(grad,ord=2)
        tol_iter[i]=tol

time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print('iterations',i)
print(OF_iter[i])
print(beta_quasi)
print(np.transpose(beta_ls_exact))
print('Tolerance=',tol)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-beta_quasi,ord=2))

('time elapsed=', 77.44199172950175)
('iterations', 99999)
84612255.8376
[-0.69455439  0.02330877 -0.03196133  0.02538503 -0.02802235  0.02159238
 -0.02782555  0.00498215  0.01533958  0.03817249 -0.0292014  -0.00445133
  0.04292589  0.03425616  0.04192465  0.01969262  0.03000818  0.05220902
  0.04888016 -0.02244421 -0.02604743 -0.0385179  -0.02231664  0.04574423

```

```

-0.04602987  0.04638269  0.04767248 -0.00495406 -0.04085386  0.01351047
-0.03731758 -0.00940024 -0.00109407 -0.02037651  0.02295996  0.05064143
 0.03090196  0.00648642  0.04892168  0.0052158  -0.00708273  0.03500935
-0.01268554  0.00202509  0.02323332  0.01934208  0.04378862  0.02574653
-0.00521638 -0.01736145  0.05041185  0.02412116 -0.02206341  0.04692197
-0.02956954 -0.0241941  -0.0450268  0.02771753  0.03523912  0.04909869
 0.04288969 -0.03328658 -0.04322499 -0.00332134  0.00306589  0.04330142
 0.04908182  0.01516352 -0.04848864  0.00754634  0.04075071  0.02911294
 0.04070567  0.04907051 -0.01998648  0.04099637  0.01332996 -0.01517363
 0.0248067  -0.04697969 -0.02677797  0.02275512  0.04800963 -0.03300447
 0.0343062  -0.04535301  0.00093947  0.04869848 -0.04790353  0.02933167
-0.02343911  0.00257443  0.03050908 -0.0114265  -0.0137986  -0.04242246
-0.04702582  0.00790206 -0.03604466  0.02715162  0.03446269]
[[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
 -2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
  4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
  4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
 -4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
 -3.8973661  -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
  2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
 -1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
 -0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
 -3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769
  4.21201842 -3.5424867  -4.50429917 -0.48099319  0.1316062  4.24485554
  4.7728776  1.38593473 -4.99300561  0.74922867  3.95908786  2.79126176
  3.99975326  4.84343426 -2.07181554  3.98489739  1.19557716 -1.65011212
  2.41928082 -4.7743059  -2.80624142  2.12898243  4.66082827 -3.45711962
  3.31431908 -4.59174474 -0.0981666  4.70477382 -4.95014286  2.87792284
 -2.45991397  0.12189604  2.88425693 -1.36533547 -1.58816079 -4.39826434
 -4.8054921  0.56007872 -3.78564008  2.50915003  3.30404243]]
('Tolerance=', 27512897.560719263)
('error=', 41.810947541672022)

```

```

In [5]: import sympy
        from pandas import DataFrame
        from sympy import *

        # For Final report we need to put 'Time elapsed', 'Iteration','Tolerance',
        ### Fill value instead of 1 2 3 .....
        t = [[215.57 , 99999 , 11.96 , 0.05617 ], #Gradient method (Armijo Rule)
              [14.09 , 99999 , 16.64 , 0.08595 ], #Gradient method
              [293.55 , 99999 , 4972171.18 , 4.24e-12 ], #Newton method (Armijo Rule)
              [192.15 , 99999 , 4972171.18 , 99.15 ], #Newton method
              [0.0035 , 2 , 2.99e-6 , 6.43e-13 ], #Quasi Newton method (Armijo Rule)
              [77.44 , 99999 , 27512897.56 , 41.81 ]] #Quasi Newton method

```



```
table = DataFrame(t, index=['Gradient method (Armijo Rule)', 'Gradient metho
                      columns=['Time elapsed', 'Iteration', 'Tolerance', 'Error
table
```

```
Out[5]:
```

	Time elapsed	Iteration	Tolerance
Gradient method (Armijo Rule)	215.5700	99999	1.196000e+01
Gradient method	14.0900	99999	1.664000e+01
Newton method (Armijo Rule)	293.5500	99999	4.972171e+06
Newton method	192.1500	99999	4.972171e+06
Quasi Newton method (Armijo Rule)	0.0035	2	2.990000e-06
Quasi Newton method	77.4400	99999	2.751290e+07

	Error
Gradient method (Armijo Rule)	5.617000e-02
Gradient method	8.595000e-02
Newton method (Armijo Rule)	4.240000e-12
Newton method	9.915000e+01
Quasi Newton method (Armijo Rule)	6.430000e-13
Quasi Newton method	4.181000e+01

1.8 E)

Estimate the value of the regression coefficients y implementing the coordinate gradient method and the stochastic gradient method. Compare their performance with the algorithms in c).

```
In [88]: #definition of partial gradient
def least_sq_reg_der_par(beta_ls,X,Y,i):
    beta_ls=np.matrix(beta_ls)
    pp=-2*(Y[i]-X[i,]*np.transpose(beta_ls))*X[i,]
    aa= np.squeeze(np.asarray(pp))
    return aa
```

1.8.1 Coordinate gradient

```
In [89]: import time
import random

(a,b)=X.shape
#alpha= 0.0000001
beta_coor=np.zeros(b) #initial value for beta
n_iter=100000 #maximim nunber iteration
OF_iter=np.zeros(n_iter)
tol_iter=np.zeros(n_iter)
alpha_iter=np.zeros(n_iter)
i=0;
tol=1000;
epsilon=1e-3;
```

```

time_start = time.clock()
while (i <= n_iter-2) and (tol>epsilon):
    i=i+1
    grad = least_sq_reg_der(beta_coor,X,Y)
    ddirect = np.zeros(b)
    j = random.randint(0, b-1)
    ddirect[j] = - grad[j]
    #####
    #      Armijo Rule-----Choose an appropriate alpha
    sigma = 0.1
    beta = 0.5
    alpha = 1
    while (least_sq_reg(beta_coor + alpha*ddirect , X, Y)> least_sq_reg(beta_coor, X, Y)):
        alpha = alpha*beta
    #####
    beta_coor = beta_coor + alpha*ddirect
    OF_iter[i] = least_sq_reg(beta_coor, X, Y)
    tol = np.linalg.norm(grad,ord=2)
    tol_iter[i] = tol
    alpha_iter[i] = alpha

time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print('iterations',i)
print(OF_iter[i])
print(beta_coor)
print(np.transpose(beta_ls_exact))
print('Tolerance=',tol)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-beta_coor,ord=2))

('time elapsed=', 163.64487752970308)
('iterations', 99999)
84919.7410451
[-2.50021396  2.18682373 -3.33330932  2.39975782 -2.92014081  2.00246299
 -2.92190745  0.38682466  1.41610014  3.80388697 -3.04994025 -0.58761597
  4.12933115  3.27038748  4.05991182  1.86812078  2.90038461  5.10213145
  4.7520726  -2.37795869 -2.74969627 -3.94575757 -2.45870403  4.49820917
 -4.75009063  4.55768429  4.65587883 -0.59582511 -4.19314339  1.24460461
 -3.89808893 -1.04057926 -0.26554226 -2.28768565  2.16739383  4.92170907
  2.93963913  0.53277756  4.79699291  0.43585696 -0.86131784  3.3522031
 -1.41222325  0.06931246  2.14531845  1.82115116  4.29541405  2.45009805
 -0.63666568 -1.91204923  4.94201387  2.30657644 -2.39688656  4.63435576
 -3.05973144 -2.51412631 -4.71816209  2.54908238  3.42649047  4.78616836
  4.2115354  -3.54344122 -4.50507367 -0.48169823  0.13078017  4.24435255
  4.7720993  1.38526205 -4.99358747  0.74918188  3.95849775  2.79061004
  3.99931219  4.84301543 -2.07213891  3.98423849  1.19489186 -1.65072046
  2.41891899 -4.77456556 -2.80681004  2.1282199  4.66004788 -3.45777118
  3.31367115 -4.59191388 -0.09910009  4.70386052 -4.95080481  2.87758379]

```

```

-2.46041758  0.12123455  2.88339249 -1.36639218 -1.58913823 -4.39890042
-4.80587062  0.55896863 -3.78642918  2.50807425  3.30328032]
[[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
-2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
 4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
 4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
-4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
-3.8973661  -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
 2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
-1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
-0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
-3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769
 4.21201842 -3.5424867  -4.50429917 -0.48099319  0.1316062  4.24485554
 4.7728776  1.38593473 -4.99300561  0.74922867  3.95908786  2.79126176
 3.99975326  4.84343426 -2.07181554  3.98489739  1.19557716 -1.65011212
 2.41928082 -4.7743059  -2.80624142  2.12898243  4.66082827 -3.45711962
 3.31431908 -4.59174474 -0.0981666  4.70477382 -4.95014286  2.87792284
-2.45991397  0.12189604  2.88425693 -1.36533547 -1.58816079 -4.39826434
-4.8054921  0.56007872 -3.78564008  2.50915003  3.30404243]]
('Tolerance=', 4.4303189821277034)
('error=', 0.0098751677873833519)

```

1.9 Stochastic gradient

```

In [90]: (a,b)=X.shape
         #alpha= 0.0000001
         beta_sto=np.zeros(b) #initial value for beta
         n_iter=100000 #maximum number iteration
         OF_iter=np.zeros(n_iter)
         tol_iter=np.zeros(n_iter)
         alpha_iter=np.zeros(n_iter)
         i=0;
         tol=1000;
         epsilon=1e-3;

         time_start = time.clock()
         while (i <= n_iter-2) and (tol>epsilon):
             i=i+1
             j = random.randint(0, b-1)
             grad_par = least_sq_reg_der_par(beta_sto,X,Y, j)
             grad = least_sq_reg_der(beta_sto,X,Y)
             ddirect = - grad_par
             #####
             #      Armijo Rule-----Choose an appropriate alpha
             sigma = 0.1
             beta = 0.5
             alpha = 1

```

```

while (least_sq_reg(beta_sto + alpha*ddirect , X, Y)> least_sq_reg(beta
    alpha = alpha*beta
#####
beta_sto = beta_sto + ddirect*alpha
OF_iter[i] = least_sq_reg(beta_sto, X, Y)
tol = np.linalg.norm(grad,ord=2)
tol_iter[i] = tol
alpha_iter[i] = alpha

time_elapsed = (time.clock() - time_start)
print('time elapsed=',time_elapsed)
print('iterations',i)
print(OF_iter[i])
print(beta_sto)
print(np.transpose(beta_ls_exact))
print('Tolerance=',tol)
print('error=',np.linalg.norm(np.transpose(beta_ls_exact)-beta_sto,ord=2)/

('time elapsed=', 322.0782561465603)
('iterations', 9999)
2113446.74323
[ 0.10969587  2.15459106 -2.73794796  2.43062366 -1.81939819  1.55858974
-2.06132185  1.17113482  0.49903371  1.57226603 -0.64545193 -0.10494674
 5.42318687  1.38397385  2.23523869  2.42252443  2.59244987  3.03017531
 4.35250852 -0.45845149 -2.30684351 -3.81494188 -1.9625614  2.64530825
-2.50652389  4.22243716  3.39618464  1.15481196 -3.14793245  2.51914073
-1.94143724 -0.90108007 -0.03820853 -1.42390156  2.09482173  1.95039804
-1.77012151  1.55584506  2.69934517  1.4177997  -1.7876631  1.38743046
-0.57572693  2.07897314  0.04608848  0.83539717  4.01770176  1.29246787
-0.06485886 -1.26403007  3.41106878  1.65507843 -1.80252834  3.02580618
-1.95996077 -0.67287535 -3.8303346  1.53249861  2.95217127  3.50176353
 3.30527248 -1.89845257 -1.97434118  0.15292702 -1.0444904  2.73142989
 2.0422938  1.27693678 -3.35879968  1.12025885  2.81241255  0.76902383
 2.01769281  0.77575851  1.53522836  2.43482166 -0.96953795 -1.15269728
 2.71898709 -2.98207682 -1.19521613  0.38061522  2.15831357 -2.96984636
 2.2403703  -0.82964194  1.08526317  4.11447167 -1.7634749  0.71495591
-1.09210182  0.74986064  2.98311974  0.95349042  1.22215274 -1.3047014
-4.55340843 -0.87160734 -2.48210947  3.10061311  2.04836621]
[[-2.82009166  2.18757262 -3.33272408  2.40050035 -2.91964431  2.00327309
-2.92131935  0.38737545  1.41670728  3.80403217 -3.04939149 -0.58693972
 4.13024247  3.27121943  4.06064646  1.86865365  2.9009471  5.1028255
 4.75282718 -2.37740208 -2.74905513 -3.94539386 -2.45767779  4.49870622
-4.74962366  4.55818668  4.65647563 -0.59536061 -4.19274504  1.24511478
-3.8973661  -1.04011885 -0.26481368 -2.28654173  2.16800687  4.92252054
 2.94043056  0.5333243  4.79756441  0.43627901 -0.86060974  3.35301358
-1.41157203  0.06996948  2.14622611  1.8217317  4.29591723  2.45076112
-0.63614187 -1.91125437  4.94259785  2.30714669 -2.39602395  4.63460325
-3.05933115 -2.51371899 -4.71726272  2.55024272  3.42703614  4.7868769

```

```

4.21201842 -3.5424867 -4.50429917 -0.48099319 0.1316062 4.24485554
4.7728776 1.38593473 -4.99300561 0.74922867 3.95908786 2.79126176
3.99975326 4.84343426 -2.07181554 3.98489739 1.19557716 -1.65011212
2.41928082 -4.7743059 -2.80624142 2.12898243 4.66082827 -3.45711962
3.31431908 -4.59174474 -0.0981666 4.70477382 -4.95014286 2.87792284
-2.45991397 0.12189604 2.88425693 -1.36533547 -1.58816079 -4.39826434
-4.8054921 0.56007872 -3.78564008 2.50915003 3.30404243]]
('Tolerance=', 259606.19719997552)
('error=', 0.74201546127562212)

```

```

In [6]: import sympy
        from pandas import DataFrame
        from sympy import *

        # For Final report we need to put 'Time elapsed', 'Iteration', 'Tolerance',
        ### Fill value instead of 1 2 3 .....
        t = [[163.64 , 99999 , 4.43 , 0.0099 ], #Coordinate gradient
              [322.08 , 99999 , 259606.20 , 0.7420 ]] #Stochastic gradient

```

```

table = DataFrame(t, index=['Coordinate gradient', 'Stochastic gradient'],

```

```

table

```

```

Out[6]:

```

	Time elapsed	Iteration	Tolerance	Error
Coordinate gradient	163.64	99999	4.43	0.0099
Stochastic gradient	322.08	99999	259606.20	0.7420

```

In [ ]:

```